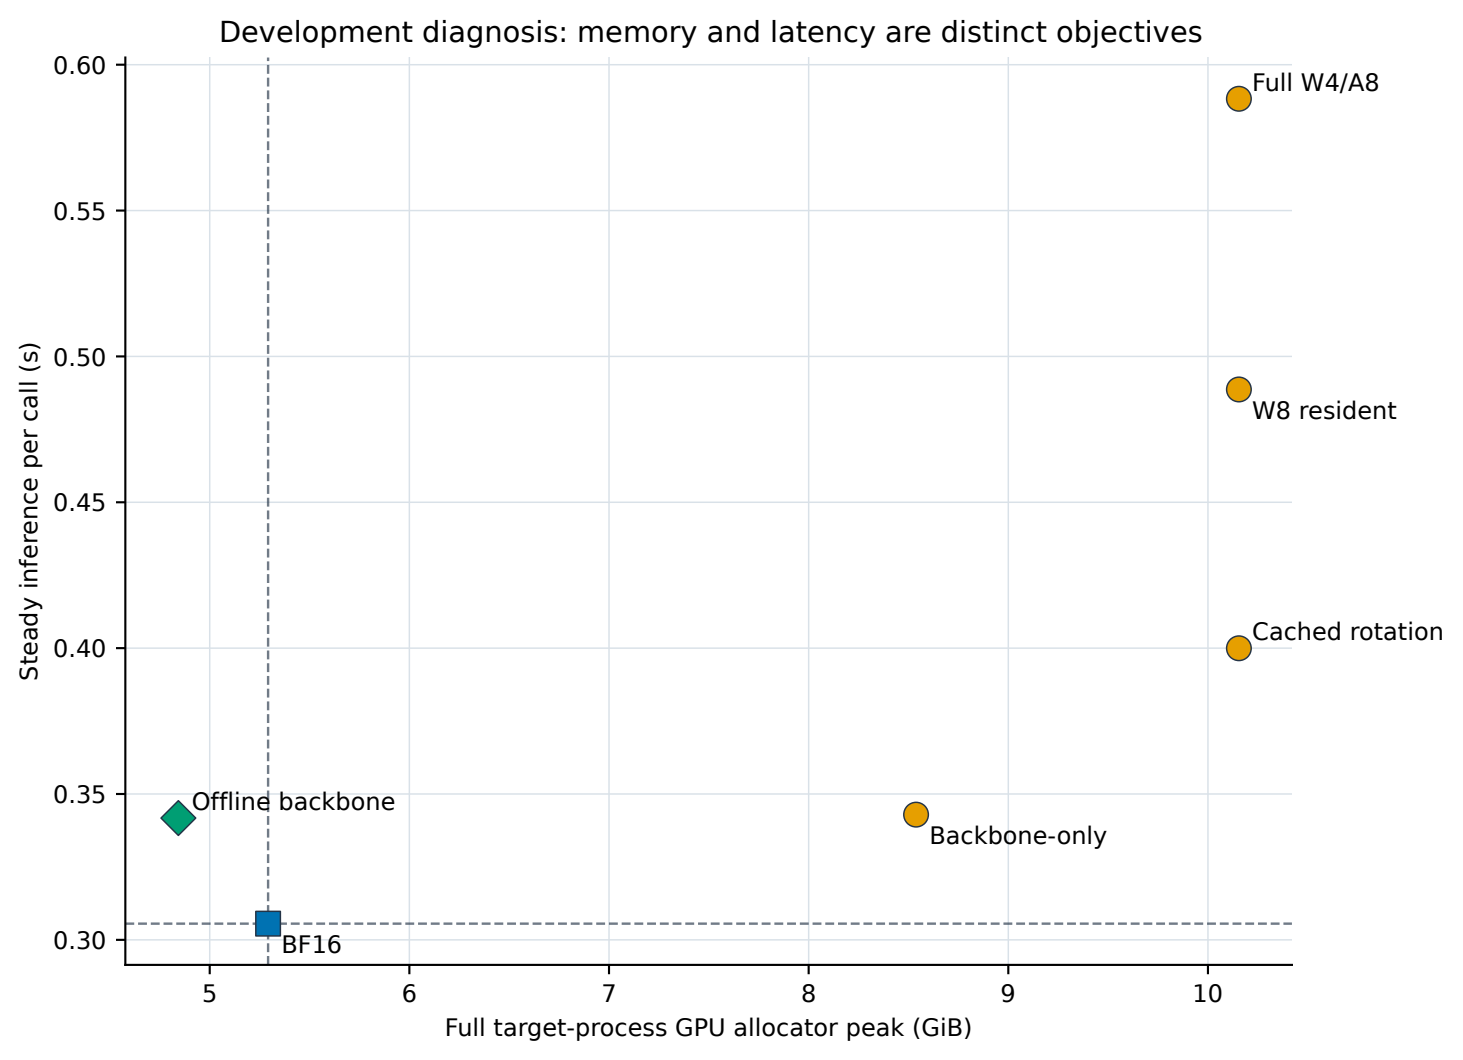
Development diagnosis: memory and latency are distinct objectives
Steady inference per call (s)
Full target-process GPU allocator peak (GiB)
0.60
0.55
0.50
0.45
0.40
0.35
0.30
5
6
7
8
9
10
Full W4/A8
W8 resident
Cached rotation
Offline backbone
Backbone-only
BF16